

New Mexico, Albuquerque: Top 5 zip codes to invest in

By: Brittney Nitta-Lee

Project Summary

The goal of this project is to analyze the historical real estate data for Albuquerque, New Mexico, and identify the top 5 zip codes for investment based on forecasted median home prices. A time series modeling approach is employed to predict future values, helping investors make informed decisions.

Business Understanding

My clients are real estate investors with a focus on properties in King County. Seeking to escape the cold weather, they are interested in exploring investment opportunities in Albuquerque, New Mexico. I will identify five potential zip codes for investment in the area and provide a list of recommendations along with suggested next steps.

Dataset

The dataset was sourced from Zillow Research and can be accessed [here](#). The dataset includes 14,723 rows, each representing a zip code, and 272 columns. The data provides median sales for every zip code from April 1996 to April 2018.

Load the Data/Filtering for Chosen Zipcodes

```
In [1]: # imports
from math import sqrt
from sklearn.metrics import mean_squared_error
import warnings
from pylab import hist, show, xticks
import itertools
from statsmodels.tsa.stattools import adfuller
from matplotlib.pylab import rcParams
import statsmodels.api as sm
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('fivethirtyeight')
warnings.filterwarnings('ignore')
import seaborn as sns
from sklearn.model_selection import TimeSeriesSplit
from statsmodels.tsa.arima.model import ARIMA
import warnings
from scipy import stats
from statsmodels.graphics.tsaplots import plot_acf
```

```
In [2]: df_zillow = pd.read_csv("zillow_data.csv")
df_zillow.head()
```

Out[2]:

	RegionID	RegionName	City	State	Metro	CountyName	SizeRank	1996-04	1996-05
0	84654	60657	Chicago	IL	Chicago	Cook	1	334200.0	335400.0
1	90668	75070	McKinney	TX	Dallas-Fort Worth	Collin	2	235700.0	236900.0
2	91982	77494	Katy	TX	Houston	Harris	3	210400.0	212200.0
3	84616	60614	Chicago	IL	Chicago	Cook	4	498100.0	500900.0
4	93144	79936	El Paso	TX	El Paso	El Paso	5	77300.0	77300.0

5 rows × 272 columns

```
In [3]: df_zillow.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14723 entries, 0 to 14722
Columns: 272 entries, RegionID to 2018-04
dtypes: float64(219), int64(49), object(4)
memory usage: 30.6+ MB
```

Data Preprocessing

The data has 272 columns and 14723 entries. For this project, I want to focus on Albuquerque, so I'll create a new dataframe called `abq_data` to include all data from Albuquerque from `Metro` column.

```
In [4]: abq_data = df_zillow.loc[df_zillow["Metro"] == "Albuquerque"]
abq_data
```

Out[4]:

	RegionID	RegionName	City	State	Metro	CountyName	SizeRank	1996-04
110	95306	87111	Albuquerque	NM	Albuquerque	Bernalillo	111	156900.0
114	95309	87114	Albuquerque	NM	Albuquerque	Bernalillo	115	139000.0
196	95314	87120	Albuquerque	NM	Albuquerque	Bernalillo	197	130900.0
268	95315	87121	Albuquerque	NM	Albuquerque	Bernalillo	269	99900.0
430	95318	87124	Rio Rancho	NM	Albuquerque	Sandoval	431	113200.0
734	95304	87109	Albuquerque	NM	Albuquerque	Bernalillo	735	136100.0
772	95307	87112	Albuquerque	NM	Albuquerque	Bernalillo	773	112200.0
893	95300	87105	South Valley	NM	Albuquerque	Bernalillo	894	96700.0
994	95317	87123	Albuquerque	NM	Albuquerque	Bernalillo	995	114400.0
1006	95305	87110	Albuquerque	NM	Albuquerque	Bernalillo	1007	110500.0

	RegionID	RegionName	City	State	Metro	CountyName	SizeRank	1996-04
1027	95303	87108	Albuquerque	NM	Albuquerque	Bernalillo	1028	96300.0
1628	95265	87031	Los Lunas	NM	Albuquerque	Valencia	1629	107300.0
1958	95321	87144	Rio Rancho	NM	Albuquerque	Sandoval	1959	135600.0
2698	95302	87107	Albuquerque	NM	Albuquerque	Bernalillo	2699	114700.0
3005	95301	87106	Albuquerque	NM	Albuquerque	Bernalillo	3006	112300.0
4333	95297	87102	Albuquerque	NM	Albuquerque	Bernalillo	4334	88200.0
5337	95239	87002	Belen	NM	Albuquerque	Valencia	5338	101200.0
6420	95316	87122	Albuquerque	NM	Albuquerque	Bernalillo	6421	232000.0
6840	95308	87113	Albuquerque	NM	Albuquerque	Bernalillo	6841	183600.0
7089	95299	87104	Albuquerque	NM	Albuquerque	Bernalillo	7090	113200.0
8082	95240	87004	Bernalillo	NM	Albuquerque	Sandoval	8083	138400.0
8163	95280	87048	Corrales	NM	Albuquerque	Sandoval	8164	193500.0
8314	95286	87059	Tijeras	NM	Albuquerque	Bernalillo	8315	142000.0
10088	95275	87043	Placitas	NM	Albuquerque	Sandoval	10089	222200.0
10207	95279	87047	Sandia Park	NM	Albuquerque	Bernalillo	10208	188700.0
11468	95274	87042	Peralta	NM	Albuquerque	Valencia	11469	111700.0
11474	95292	87068	Bosque Farms	NM	Albuquerque	Valencia	11475	122200.0
12222	95244	87008	Cedar Crest	NM	Albuquerque	Bernalillo	12223	186000.0

28 rows × 272 columns

The data contains 28 different zipcodes in Albuquerque. I will drop `RegionID` and `SizeRank` as I don't need those columns for this project. I will also look for any missing data.

```
In [5]: abq_data.drop(['RegionID', 'SizeRank'], axis=1, inplace=True)
```

```
In [6]: # checking for null values
abq_data[pd.isna(abq_data).any(axis=1)]
```

```
Out[6]:
```

RegionName	City	State	Metro	CountyName	1996-04	1996-05	1996-06	1996-07	1996-08	...	2017-07

0 rows × 270 columns

```
In [7]: abq_data['CountyName'].unique()
```

```
Out[7]: array(['Bernalillo', 'Sandoval', 'Valencia'], dtype=object)
```

There's three different counties listed in `abq_data` . Next, I'll explore the data and narrow down the zipcodes I want to use for my time series model.

```
In [8]: abq_data2 = abq_data.copy()
```

Exploratory Data Analysis

To provide more insightful information to our stakeholders, I will create a new column ROI in the `abq_data`. This column will represent the total Return on Investment (ROI) for each zip code, which is a more comprehensive measure of profitability. I'll also include the cumulative percent change as these are features to evaluate the performance of my model.

```
In [9]: # create metrics to judge return on investment (ROI)

# total ROI
abq_data['Total_ROI'] = (abq_data['2018-04'] -
                        abq_data['1996-04']) / abq_data['1996-04']

# cumulative percent change
abq_data['Cumulative_Percent_Change'] = (abq_data['2018-04'] / abq_data['1996-04'] - 1) * 100

# re-assign to `abq_roi` dataframe with required columns
abq_roi = abq_data.loc[:, ['RegionName', 'City', 'State', 'Metro', 'CountyName', 'Total_ROI', 'Cumulative_Percent_Change']]
abq_roi.head()
```

Out[9]:

	RegionName	City	State	Metro	CountyName	Total_ROI	Cumulative_Percent_Change
110	87111	Albuquerque	NM	Albuquerque	Bernalillo	0.695347	169.1
114	87114	Albuquerque	NM	Albuquerque	Bernalillo	0.565468	156.1
196	87120	Albuquerque	NM	Albuquerque	Bernalillo	0.486631	148.1
268	87121	Albuquerque	NM	Albuquerque	Bernalillo	0.434434	143.1
430	87124	Rio Rancho	NM	Albuquerque	Sandoval	0.595406	159.1

```
In [10]: # Find the top 10 zipcodes in Albuquerque
abq_roi.sort_values('Total_ROI', ascending=False).head(10)
```

Out[10]:

	RegionName	City	State	Metro	CountyName	Total_ROI	Cumulative_Percent_Change
8163	87048	Corrales	NM	Albuquerque	Sandoval	1.204651	204.1
6420	87122	Albuquerque	NM	Albuquerque	Bernalillo	1.180603	201.1
3005	87106	Albuquerque	NM	Albuquerque	Bernalillo	1.102404	201.1
10088	87043	Placitas	NM	Albuquerque	Sandoval	0.903690	199.1
7089	87104	Albuquerque	NM	Albuquerque	Bernalillo	0.750883	179.1
734	87109	Albuquerque	NM	Albuquerque	Bernalillo	0.713446	171.1
8314	87059	Tijeras	NM	Albuquerque	Bernalillo	0.702817	171.1
110	87111	Albuquerque	NM	Albuquerque	Bernalillo	0.695347	169.1

	RegionName	City	State	Metro	CountyName	Total_ROI	Cumulative_Percen
2698	87107	Albuquerque	NM	Albuquerque	Bernalillo	0.684394	16
1027	87108	Albuquerque	NM	Albuquerque	Bernalillo	0.666667	16

New Mexico's largest county is Bernalillo, which contains 47 zipcodes. In comparison, Sandoval county only contains 18 zipcodes. Interestingly, despite its smaller size, the median household income in Sandoval county is \$10,000 higher than that of Bernalillo county.

In [11]: `abq_data.sort_values`

```
Out[11]: <bound method DataFrame.sort_values of
Metro CountyName 1996-04 \
110      87111 Albuquerque NM Albuquerque Bernalillo 156900.0
114      87114 Albuquerque NM Albuquerque Bernalillo 139000.0
196      87120 Albuquerque NM Albuquerque Bernalillo 130900.0
268      87121 Albuquerque NM Albuquerque Bernalillo 99900.0
430      87124 Rio Rancho NM Albuquerque Sandoval 113200.0
734      87109 Albuquerque NM Albuquerque Bernalillo 136100.0
772      87112 Albuquerque NM Albuquerque Bernalillo 112200.0
893      87105 South Valley NM Albuquerque Bernalillo 96700.0
994      87123 Albuquerque NM Albuquerque Bernalillo 114400.0
1006     87110 Albuquerque NM Albuquerque Bernalillo 110500.0
1027     87108 Albuquerque NM Albuquerque Bernalillo 96300.0
1628     87031 Los Lunas NM Albuquerque Valencia 107300.0
1958     87144 Rio Rancho NM Albuquerque Sandoval 135600.0
2698     87107 Albuquerque NM Albuquerque Bernalillo 114700.0
3005     87106 Albuquerque NM Albuquerque Bernalillo 112300.0
4333     87102 Albuquerque NM Albuquerque Bernalillo 88200.0
5337     87002 Belen NM Albuquerque Valencia 101200.0
6420     87122 Albuquerque NM Albuquerque Bernalillo 232000.0
6840     87113 Albuquerque NM Albuquerque Bernalillo 183600.0
7089     87104 Albuquerque NM Albuquerque Bernalillo 113200.0
8082     87004 Bernalillo NM Albuquerque Sandoval 138400.0
8163     87048 Corrales NM Albuquerque Sandoval 193500.0
8314     87059 Tijeras NM Albuquerque Bernalillo 142000.0
10088    87043 Placitas NM Albuquerque Sandoval 222200.0
10207    87047 Sandia Park NM Albuquerque Bernalillo 188700.0
11468    87042 Peralta NM Albuquerque Valencia 111700.0
11474    87068 Bosque Farms NM Albuquerque Valencia 122200.0
12222    87008 Cedar Crest NM Albuquerque Bernalillo 186000.0

1996-05 1996-06 1996-07 1996-08 ... 2017-09 2017-10 2017-11 \
110 155600.0 154200.0 152800.0 151400.0 ... 253700 254200 256000
114 139200.0 139300.0 139400.0 139500.0 ... 209600 210700 211800
196 130500.0 130000.0 129500.0 129000.0 ... 191400 192300 193000
268 99600.0 99600.0 99600.0 99700.0 ... 139800 140100 140600
430 113400.0 113600.0 113600.0 113700.0 ... 175300 176700 178800
734 135800.0 135500.0 135300.0 135200.0 ... 229100 228700 229000
772 112400.0 112600.0 112800.0 113100.0 ... 170900 171000 172100
893 95800.0 95100.0 94600.0 94300.0 ... 136200 136900 137600
994 114700.0 115000.0 115200.0 115500.0 ... 181000 181700 181800
1006 110600.0 110700.0 110900.0 111100.0 ... 180400 180300 180500
1027 96600.0 97000.0 97500.0 97900.0 ... 155500 155400 156200
1628 107200.0 107100.0 106900.0 106800.0 ... 135900 137000 138100
1958 135300.0 135000.0 134700.0 134500.0 ... 183700 184000 185400
2698 115000.0 115400.0 115800.0 116300.0 ... 189300 190000 190100
3005 112300.0 112400.0 112400.0 112400.0 ... 221000 221400 223700
4333 88000.0 87900.0 87900.0 87900.0 ... 126200 128700 129500
5337 101300.0 101200.0 101200.0 101200.0 ... 107100 107900 109500
```

6420	232000.0	232000.0	232200.0	232400.0	...	497500	499600	503100
6840	183900.0	184200.0	184500.0	184800.0	...	269200	269400	270300
7089	113300.0	113400.0	113500.0	113600.0	...	181600	184100	186500
8082	138400.0	138400.0	138400.0	138400.0	...	203500	203800	205600
8163	193700.0	193800.0	194000.0	194200.0	...	408900	408300	411400
8314	141900.0	141700.0	141500.0	141300.0	...	232300	231900	232900
10088	222400.0	222700.0	223100.0	223600.0	...	392600	395000	402400
10207	188800.0	188800.0	188800.0	188700.0	...	275000	274800	277200
11468	112000.0	112100.0	112200.0	112300.0	...	166400	167000	167200
11474	122600.0	123000.0	123400.0	123800.0	...	193800	195900	196100
12222	185600.0	185100.0	184700.0	184200.0	...	277300	279000	278800

	2017-12	2018-01	2018-02	2018-03	2018-04	Total_ROI	\
110	258600	260500	262300	264500	266000	0.695347	
114	213100	214100	215200	216600	217600	0.565468	
196	193400	193300	193600	194300	194600	0.486631	
268	141100	141200	141900	142800	143300	0.434434	
430	180700	181400	181000	180500	180600	0.595406	
734	230100	231100	231800	232700	233200	0.713446	
772	173500	174400	176000	178500	180100	0.605169	
893	138200	138400	138500	138600	138400	0.431231	
994	181600	181200	182200	184300	185300	0.619755	
1006	180500	179900	180400	182400	184000	0.665158	
1027	157300	158000	158600	159700	160500	0.666667	
1628	139200	139900	141100	142400	143200	0.334576	
1958	186700	187100	187100	187600	188100	0.387168	
2698	189700	189600	190000	191600	193200	0.684394	
3005	227500	230500	232700	234700	236100	1.102404	
4333	129300	129400	129600	130700	132200	0.498866	
5337	111100	111800	112900	114300	115100	0.137352	
6420	505700	508200	509300	508200	505900	1.180603	
6840	271300	271500	272900	274600	274600	0.495643	
7089	188800	190300	192100	195300	198200	0.750883	
8082	207500	207600	207100	209500	212800	0.537572	
8163	414400	415000	418200	423200	426600	1.204651	
8314	234200	234900	237200	240300	241800	0.702817	
10088	410900	416700	419400	421000	423000	0.903690	
10207	282200	286300	290500	295100	298200	0.580286	
11468	167700	168600	170200	171500	172000	0.539839	
11474	195700	196900	198600	199700	200400	0.639935	
12222	278200	279200	280800	281900	282000	0.516129	

	Cumulative_Percent_Change
110	169.534736
114	156.546763
196	148.663102
268	143.443443
430	159.540636
734	171.344600
772	160.516934
893	143.123061
994	161.975524
1006	166.515837
1027	166.666667
1628	133.457596
1958	138.716814
2698	168.439407
3005	210.240427
4333	149.886621
5337	113.735178
6420	218.060345
6840	149.564270
7089	175.088339
8082	153.757225
8163	220.465116

```

8314          170.281690
10088         190.369037
10207         158.028617
11468         153.983885
11474         163.993453
12222         151.612903

```

```
[28 rows x 272 columns]>
```

Visualization

Now that I have selected the top 10 zipcodes in Albuquerque, I'll create a bar graph to see the growth percentage.

```


In [12]: # formatting for the rest of our visualizations
font = {'family': 'DejaVu Sans',
        'weight': 'bold',
        'size': 22}

plt.rc('font', **font)

In [13]: # Sort values by Total_ROI
abq_roi_sorted = abq_roi.sort_values(by='Total_ROI', ascending=False)
ax = abq_roi_sorted.head(10).plot.bar(x='RegionName', y='Total_ROI', figsize=(
    16, 8), alpha=0.5, edgecolor="black", linewidth=2)
plt.title('Top 10 Zip Codes in Albuquerque', fontsize=25)
plt.legend('')
ax.set_yticklabels(["0%", "20%", "40%", "60%", "80%", "100%"])
plt.xlabel('Zip Codes', fontsize=20)
plt.ylabel('Growth Percentage', fontsize=20)

```

```
Out[13]: Text(0, 0.5, 'Growth Percentage')
```



```

In [14]: # Create new dataframe for top 10 zipcodes in Albuquerque
abq_10 = abq_roi.sort_values(by='Total_ROI', ascending=False).head(10)
abq_10

```

Out[14]:

	RegionName	City	State	Metro	CountyName	Total_ROI	Cumulative_Percent
8163	87048	Corrales	NM	Albuquerque	Sandoval	1.204651	2
6420	87122	Albuquerque	NM	Albuquerque	Bernalillo	1.180603	2
3005	87106	Albuquerque	NM	Albuquerque	Bernalillo	1.102404	2
10088	87043	Placitas	NM	Albuquerque	Sandoval	0.903690	19
7089	87104	Albuquerque	NM	Albuquerque	Bernalillo	0.750883	17
734	87109	Albuquerque	NM	Albuquerque	Bernalillo	0.713446	17
8314	87059	Tijeras	NM	Albuquerque	Bernalillo	0.702817	17
110	87111	Albuquerque	NM	Albuquerque	Bernalillo	0.695347	16
2698	87107	Albuquerque	NM	Albuquerque	Bernalillo	0.684394	16
1027	87108	Albuquerque	NM	Albuquerque	Bernalillo	0.666667	16

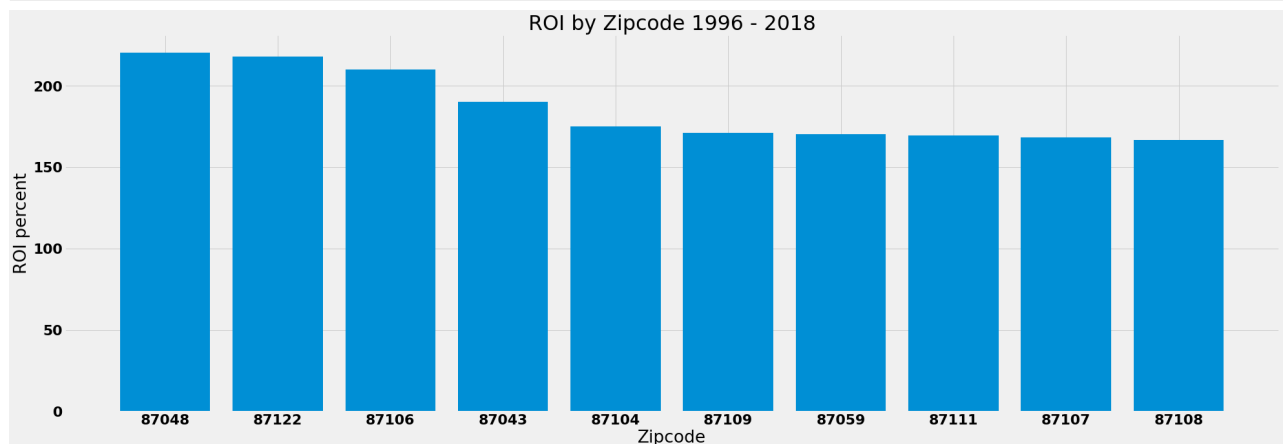
In [15]:

```

abq_10['RegionName'] = abq_10['RegionName'].astype(str)

# Plotting the historical data
fig, ax = plt.subplots(figsize=(30,10))
plt.bar(abq_10.RegionName, abq_10['Cumulative_Percent_Change'])
plt.title('ROI by Zipcode 1996 - 2018')
plt.xlabel('Zipcode')
plt.ylabel('ROI percent')
plt.show()

```



This graph shows the top 10 zipcodes with the highest ROI over the period of 1996 to 2018. There's subtle changes between 87101 and 87108.

In [16]:

```

# list of top 5 zipcodes
region_list = ['87048', '87122', '87106', '87043', '87104']

# filter rows with desired RegionNames
abq_top_5 = abq_data[abq_data['RegionName'].isin(region_list)]

abq_top_5 = abq_top_5.drop(columns=['Total_ROI', 'Cumulative_Percent_Change'])

# display new dataframe
abq_top_5.head()

```


Out[16]:

	RegionName	City	State	Metro	CountyName	1996-04	1996-05	1996-06
3005	87106	Albuquerque	NM	Albuquerque	Bernalillo	112300.0	112300.0	112400.0
6420	87122	Albuquerque	NM	Albuquerque	Bernalillo	232000.0	232000.0	232000.0
7089	87104	Albuquerque	NM	Albuquerque	Bernalillo	113200.0	113300.0	113400.0
8163	87048	Corrales	NM	Albuquerque	Sandoval	193500.0	193700.0	193800.0
10088	87043	Placitas	NM	Albuquerque	Sandoval	222200.0	222400.0	222700.0

5 rows × 270 columns

In [17]: abq_top_5

Out[17]:

	RegionName	City	State	Metro	CountyName	1996-04	1996-05	1996-06
3005	87106	Albuquerque	NM	Albuquerque	Bernalillo	112300.0	112300.0	112400.0
6420	87122	Albuquerque	NM	Albuquerque	Bernalillo	232000.0	232000.0	232000.0
7089	87104	Albuquerque	NM	Albuquerque	Bernalillo	113200.0	113300.0	113400.0
8163	87048	Corrales	NM	Albuquerque	Sandoval	193500.0	193700.0	193800.0
10088	87043	Placitas	NM	Albuquerque	Sandoval	222200.0	222400.0	222700.0

5 rows × 270 columns

Reshape from Wide to Long Format

Now that I have the data that I want to use, I'll reshape the data from a wideformat to a long format. The `abq_data` contains the wide format dataset.

```
In [18]: def melt_data(df):
# Melt data into wide version
melted = pd.melt(df, id_vars=['RegionName', 'City', 'State', 'Metro', 'Count
# Create new column as datetime variable
melted['time'] = pd.to_datetime(melted['time'], infer_datetime_format=True)
# Remove rows with missing values
melted = melted.dropna(subset=['value'])
# set `time` as index
melted.set_index('time', inplace=True)
return melted.groupby('time').aggregate({'value': 'mean'})
```

```
In [19]: # Create new data frame for zipcodes
abq_zip = [zip_code for zip_code in abq_10['RegionName']]
```

In [20]: abq_zip

```
Out[20]: ['87048',
          '87122',
          '87106',
          '87043',
          '87104',
          '87109',
          '87059',
          '87111',
          '87107',
          '87108']
```

Final top 10 zipcodes

I want to create a visualization for the top 10 zipcodes. The visualization helps with detecting and trends or patterns.

```
In [21]: abq_df = pd.DataFrame()
for i in abq_top_5['RegionName']:
    x = melt_data(abq_top_5[abq_top_5['RegionName'] == i])
    abq_df = pd.concat([abq_df, x], axis=1)
    abq_df.rename(columns = {'value':i}, inplace = True)

# Display results
abq_df.head(10)
```

```
Out[21]:
```


	87106	87122	87104	87048	87043
time					
1996-04-01	112300.0	232000.0	113200.0	193500.0	222200.0
1996-05-01	112300.0	232000.0	113300.0	193700.0	222400.0
1996-06-01	112400.0	232000.0	113400.0	193800.0	222700.0
1996-07-01	112400.0	232200.0	113500.0	194000.0	223100.0
1996-08-01	112400.0	232400.0	113600.0	194200.0	223600.0
1996-09-01	112400.0	232800.0	113700.0	194400.0	224100.0
1996-10-01	112500.0	233200.0	113800.0	194700.0	224700.0
1996-11-01	112600.0	233800.0	114000.0	194700.0	225000.0
1996-12-01	112800.0	234500.0	114100.0	194600.0	224900.0
1997-01-01	113100.0	235300.0	114300.0	194700.0	225100.0

To make it easier to analyze and visualize the data, I'm going to resample the time series to a monthly frequency with the start of each month as the observation point. This will aggregate the data for each month and create a new dataframe with the data for the first day of each month. This is particularly useful for forecasting and filling in any gaps in the data.

```
In [22]: abq_df = abq_df.asfreq('MS')
```

```
In [23]: abq_df.plot(figsize=(15, 7))
plt.xlabel('Year')
plt.ylabel('Value')
```

```
plt.title('Median Home Value: 1996-2018')
plt.show();
```



Now we can see our training data and the testing data. 80% of our data is in the train and 20% of the data is in test. It looks like the split point is in 2014. The trend line has an upward pattern up until 2008 where you can see a downward trend. This is due to the Great Recession. The "Subprime Mortgage Crisis" was a period of time (2007 to 2010) when there was an increase in the number of high-risk mortgages that went into default and caused a ripple effect on the housing market and broader economy. This is important to highlight as I am not including this data into my modeling.

[More information about the Subprime Mortgage Crisis](<https://www.history.com/topics/21st-century/recession>)

Train Validation Split and Naive Model

I'll create a naive model by shifting the `train` data by one day to create a simple baseline model for comparison with my SARIMAX model. The naive model predicts that the current value is the same as the value from the previous day and does not take into account any patterns or trends in the data. The purpose of this model is to establish a baseline performance metric, which is the root mean squared error (RMSE).

```
In [24]: # Create new dataframe with datapoints beginning in 2011
abq_df_new = abq_df['2011-01-01':]
```

```
In [25]: # Get a list of unique zipcodes (column names)
unique_zipcodes = abq_df_new.columns

# Initialize an empty dictionary to store the RMSE values for each zipcode
rmse_dict = {}

# Define the train-test split ratio
split_ratio = 0.8

# Loop through the unique zipcodes
```

```

for zipcode in unique_zipcodes:
    # Get the data for the current zipcode
    data = abq_df_new[zipcode]

    # Calculate the index for the train-test split
    cutoff = int(len(data) * split_ratio)

    # Split the data into train and test sets
    train = data[:cutoff]
    test = data[cutoff:]

    # Shift the train data by 1 time step to create the naive model predictions
    naive_predictions = train.shift(1)

    # Calculate the RMSE between the actual values and the naive model predictions
    rmse_naive_train = np.sqrt(np.mean((train[1:] - naive_predictions[1:])**2))

    # Shift the test data by 1 time step to create the naive model predictions
    naive_predictions_test = test.shift(1)

    # Calculate the RMSE between the actual values and the naive model predictions
    rmse_naive_test = np.sqrt(np.mean((test[1:] - naive_predictions_test[1:])**2))

    # Add the RMSE values to the dictionary with the zipcode as the key
    rmse_dict[zipcode] = {'train': rmse_naive_train, 'test': rmse_naive_test}

# Print the baseline RMSE values for each zipcode
print('Baseline RMSE values for each zipcode:')
for zipcode, rmse_values in rmse_dict.items():
    print(f'Zipcode {zipcode}: Train RMSE = {rmse_values["train"]:.2f}, Test RMSE = {rmse_values["test"]:.2f}')

```

Baseline RMSE values for each zipcode:

```

Zipcode 87106: Train RMSE = 1992.74, Test RMSE = 1763.69
Zipcode 87122: Train RMSE = 2483.10, Test RMSE = 2151.74
Zipcode 87104: Train RMSE = 1084.01, Test RMSE = 1848.21
Zipcode 87048: Train RMSE = 3602.07, Test RMSE = 2825.51
Zipcode 87043: Train RMSE = 2222.06, Test RMSE = 3517.35

```

Now that I have my baseline RMSE for each zipcode, I'll use the RMSE as a baseline to evaluate my models.

SARIMAX Modeling

Now that I have my baseline model, I will use a SARIMA model to forecast the median value of the top 5 zipcodes in Albuquerque. The top five zipcodes will have individual predictions and forecast results. I will evaluate each model by calculate the Root Mean Squared Error (RMSE). I will also include the 95% confidence interval, which will give a range of values within the true future value is likely to fall, with of course, 95% confidence.

I chose a SARIMAX model for forecasting the top 5 zipcodes for multiple reasons:

1. Seasonality - Real Estate prices exhibit seasonal patterns.
2. Autoregressive and Moving Average Components - SARIMAX is an extension of the ARIMA model, which combines AR and moving average components.
3. Flexibility - Specifiy different orders for AR, MA, Seasonal, which allows for fine-tuning the model to better fit the data for each zipcode.

4. Interpretability - Interpretable results.

Albuquerque, New Mexico 87106

```
In [26]: # Define train and validation datasets based 20% and 80%
training_data = abq_df_new[87106][:cutoff]
validation_data = abq_df_new[87106][cutoff:]

# Define the range of parameters for p, d, q, P, D, Q, and s
p = d = q = range(0, 2)
P = D = Q = range(0, 2)
s = 12 # monthly data

# Generate a list of all possible combinations of parameters
pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], s) for x in itertools.product(P, D, Q)]

# Initialize variables to store the best parameters and the lowest RMSE
best_params = (0, 0, 0, 0, 0, 0, 0)
lowest_rmse = float('inf')

# Loop through all possible combinations of parameters
for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            # Fit a SARIMAX model with the current combination of parameters
            model = sm.tsa.statespace.SARIMAX(training_data,
                                                order=param,
                                                seasonal_order=param_seasonal,
                                                enforce_stationarity=False,
                                                enforce_invertibility=False).fit()

            # Make predictions on the validation data
            predictions = model.predict(start=validation_data.index[0], end=validation_data.index[-1])

            # Calculate the RMSE of the predictions
            rmse = np.sqrt(np.mean((predictions - validation_data)**2))

            # Update the best parameters and lowest RMSE if the current RMSE is
            if rmse < lowest_rmse:
                best_params = (param, param_seasonal)
                lowest_rmse = rmse

        except ValueError: # skip combinations that fail to converge or produce NaNs
            continue

print(f'Best parameters: {best_params}')
print(f'Lowest RMSE: {lowest_rmse:.2f}')
```

```
Best parameters: ((1, 0, 0), (0, 0, 0, 12))
Lowest RMSE: 4536.00
```

The root mean squared error (RMSE) is \$4,536, which represents the average difference between the actual data points and the predictions made by the SARIMAX model. This is higher than our baseline model and it indicates that the model may not be adequately capturing the data's structure.

```
In [27]: #Define SARIMAX model and fit data save as sarima_mod1
sarima_mod1 = sm.tsa.statespace.SARIMAX(abq_df_new[87106],
                                         order=(1, 0, 0),
                                         seasonal_order=(0, 0, 0, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False).fit()

# Forecast 52 months into the future (4 years)
forecast1 = sarima_mod1.get_forecast(steps=52).summary_frame()

# Calculate the mean of the last week of the forecast as the predicted value
forecast1_mean = round(forecast1['mean'][51])


#Calculate the difference between lower and upper 95% confidence intervals of th
low_int1 = round(forecast1['mean_ci_lower'][51])
high_int1 = round(forecast1['mean_ci_upper'][51])

#Calculate the difference between the upper and lower confidence intervals
ci_delta1 = round(high_int1 - low_int1)

# Print predicted value and confidence intervals
print(f'Albuquerque, NM, 87106:')
print(f'95% confidence: ${low_int1} and ${high_int1}')
print(f'Confidence range: ${ci_delta1}')

# Plot the original data and predicted values with confidence intervals
fig, ax = plt.subplots(figsize=(15, 7))
plt.plot(abq_df_new[87106])
plt.plot(forecast1['mean'])
ax.fill_between(forecast1.index, forecast1['mean_ci_lower'],
               forecast1['mean_ci_upper'], color='k', alpha=0.1)
plt.title('Albuquerque, NM, 87106')
plt.legend(['Original', 'Predicted'], loc='lower right')
plt.xlabel('Year')
plt.ylabel('Median Home Price')
plt.show()
```

Albuquerque, NM, 87106:
 95% confidence: \$236629.0 and \$293381.0
 Confidence range: \$56752.0



The result shows the forecast for the median home price in zipcode 87106. The 95% confidence interval is a measure of uncertainty around the predicted value. The 95% confidence interval is between 236,629 and 293,381. This means that based on the model's predictions, there is a 95% probability that the true median home price will fall within this range. The confidence range is \$56,752 which is the difference between the upper and lower bounds of the confidence interval.

Albuquerque, New Mexico 87122

```
In [28]: # Define train and validation datasets based 20% and 80%
training_data2 = abq_df_new[87122][:cutoff]
validation_data2 = abq_df_new[87122][cutoff:]

# Define the range of parameters for p, d, q, P, D, Q, and s
p = d = q = range(0, 2)
P = D = Q = range(0, 2)
s = 12 # monthly data

# Generate a list of all possible combinations of parameters
pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], s) for x in itertools.product(P, D, Q)]

# Initialize variables to store the best parameters and the lowest RMSE
best_params = (0, 0, 0, 0, 0, 0, 0)
lowest_rmse = float('inf')

# Loop through all possible combinations of parameters
for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            # Fit a SARIMAX model with the current combination of parameters
            model2 = sm.tsa.statespace.SARIMAX(training_data2,
                                                order=param,
                                                seasonal_order=param_seasonal,
                                                enforce_stationarity=False,
                                                enforce_invertibility=False).fit()

            # Make predictions on the validation data
            predictions2 = model2.predict(start=validation_data2.index[0], end=v

            # Calculate the RMSE of the predictions
            rmse2 = np.sqrt(np.mean((predictions2 - validation_data2)**2))

            # Update the best parameters and lowest RMSE if the current RMSE is
            if rmse2 < lowest_rmse:
                best_params = (param, param_seasonal)
                lowest_rmse = rmse2

        except ValueError: # skip combinations that fail to converge or produce
            continue

print(f'Best parameters: {best_params}')
print(f'Lowest RMSE: {lowest_rmse:.2f}')
```

```
Best parameters: ((0, 1, 1), (0, 0, 1, 12))
Lowest RMSE: 3155.22
```

The RMSE is a little closer to our baseline model. Due to the computational time, I'll leave the parameters to 0,2.

```
In [29]: #Define SARIMAX model and fit data save as sarima_mod1
sarima_mod2 = sm.tsa.statespace.SARIMAX(abq_df_new[87122],
                                         order=(0, 1, 1),
                                         seasonal_order=(0, 0, 1, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False).fit()

# # Forecast 52 months into the future (4 years)
forecast2 = sarima_mod2.get_forecast(steps=52).summary_frame()

# Calculate the mean of the last week of the forecast as the predicted value
forecast2_mean = round(forecast2['mean'][51])


#Calculate the difference between lower and upper 95% confidence intervals of the
low_int2 = round(forecast2['mean_ci_lower'][51])
high_int2 = round(forecast2['mean_ci_upper'][51])

#Calculate the difference between the upper and lower confidence intervals
ci_delta2 = round(high_int2 - low_int2)

# Print predicted value and confidence intervals
print(f'Albuquerque, NM, 87122:')
print(f'95% confidence: ${low_int1} and ${high_int1}')
print(f'Confidence range: ${ci_delta1}')

# Plot the original data and predicted values with confidence intervals
fig, ax = plt.subplots(figsize=(15, 7))
plt.plot(abq_df_new[87122])
plt.plot(forecast2['mean'])
ax.fill_between(forecast2.index, forecast2['mean_ci_lower'],
               forecast2['mean_ci_upper'], color='k', alpha=0.1)
plt.title('Albuquerque, NM, 87122')
plt.legend(['Original', 'Predicted'], loc='lower right')
plt.xlabel('Year')
plt.ylabel('Median Home Price')
plt.show()
```

Albuquerque, NM, 87122:
 95% confidence: \$236629.0 and \$293381.0
 Confidence range: \$56752.0



Zipcode 87122 has the same results as 87106. Let's move on and see the results of our other zipcodes.

Albuquerque, NM, 87104

```
In [30]: # Define train and validation datasets based 20% and 80%
training_data3 = abq_df_new[87104][:cutoff]
validation_data3 = abq_df_new[87104][cutoff:]

# Define the range of parameters for p, d, q, P, D, Q, and s
p = d = q = range(0, 2)
P = D = Q = range(0, 2)
s = 12 # monthly data

# Generate a list of all possible combinations of parameters
pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], s) for x in itertools.product(P, D, Q)]

# Initialize variables to store the best parameters and the lowest RMSE
best_params = (0, 0, 0, 0, 0, 0, 0)
lowest_rmse = float('inf')

# Loop through all possible combinations of parameters
for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            # Fit a SARIMAX model with the current combination of parameters
            model3 = sm.tsa.statespace.SARIMAX(training_data3,
                                                order=param,
                                                seasonal_order=param_seasonal,
                                                enforce_stationarity=False,
                                                enforce_invertibility=False).fit()

            # Make predictions on the validation data
            predictions3 = model3.predict(start=validation_data3.index[0], end=v

            # Calculate the RMSE of the predictions
            rmse3 = np.sqrt(np.mean((predictions3 - validation_data3)**2))
```

```

        # Update the best parameters and lowest RMSE if the current RMSE is
        if rmse3 < lowest_rmse:
            best_params = (param, param_seasonal)
            lowest_rmse = rmse3

    except ValueError: # skip combinations that fail to converge or produce
        continue

print(f'Best parameters: {best_params}')
print(f'Lowest RMSE: {lowest_rmse:.2f}')

```

Best parameters: ((1, 0, 1), (1, 0, 1, 12))
 Lowest RMSE: 5235.02

The RMSE is a lot higher than our baseline model.

```

In [31]: #Define SARIMAX model and fit data save as sarima_mod1
sarima_mod3 = sm.tsa.statespace.SARIMAX(abq_df_new[87104],
                                         order=(1, 0, 1),
                                         seasonal_order=(1, 0, 1, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False).fit()

# Forecast 52 months into the future (4 years)
forecast3 = sarima_mod3.get_forecast(steps=52).summary_frame()

# Calculate the mean of the last week of the forecast as the predicted value
forecast3_mean = round(forecast3['mean'][51])

#Calculate the difference between lower and upper 95% confidence intervals of th
low_int3 = round(forecast3['mean_ci_lower'][51])
high_int3 = round(forecast3['mean_ci_upper'][51])


#Calculate the difference between the upper and lower confidence intervals
ci_delta3 = round(high_int3 - low_int3)

# Print predicted value and confidence intervals
print(f'Albuquerque, NM, 87104:')
print(f'95% confidence: ${low_int1} and ${high_int1}')
print(f'Confidence range: ${ci_delta1}')

# Plot the original data and predicted values with confidence intervals
fig, ax = plt.subplots(figsize=(15, 7))
plt.plot(abq_df_new[87104])
plt.plot(forecast3['mean'])
ax.fill_between(forecast3.index, forecast3['mean_ci_lower'],
               forecast3['mean_ci_upper'], color='k', alpha=0.1)
plt.title('Albuquerque, NM, 87104')
plt.legend(['Original', 'Predicted'], loc='lower right')
plt.xlabel('Year')
plt.ylabel('Median Home Price')
plt.show()

```

Albuquerque, NM, 87104:
 95% confidence: \$236629.0 and \$293381.0
 Confidence range: \$56752.0



This also gave the same results as the previous zipcodes.

Albuquerque, NM, 87048

```
In [32]: # Define train and validation datasets based 20% and 80%
training_data4 = abq_df_new[87048][:cutoff]
validation_data4 = abq_df_new[87048][cutoff:]

# Define the range of parameters for p, d, q, P, D, Q, and s
p = d = q = range(0, 2)
P = D = Q = range(0, 2)
s = 12 # monthly data

# Generate a list of all possible combinations of parameters
pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], s) for x in itertools.product(P, D, Q)]

# Initialize variables to store the best parameters and the lowest RMSE
best_params = (0, 0, 0, 0, 0, 0, 0)
lowest_rmse = float('inf')

# Loop through all possible combinations of parameters
for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            # Fit a SARIMAX model with the current combination of parameters
            model4 = sm.tsa.statespace.SARIMAX(training_data4,
                                                order=param,
                                                seasonal_order=param_seasonal,
                                                enforce_stationarity=False,
                                                enforce_invertibility=False).fit()

            # Make predictions on the validation data
            predictions4 = model4.predict(start=validation_data4.index[0], end=v

            # Calculate the RMSE of the predictions
            rmse4 = np.sqrt(np.mean((predictions4 - validation_data4)**2))

            # Update the best parameters and lowest RMSE if the current RMSE is
```

```

        if rmse4 < lowest_rmse:
            best_params = (param, param_seasonal)
            lowest_rmse = rmse4

    except ValueError: # skip combinations that fail to converge or produce
        continue

print(f'Best parameters: {best_params}')
print(f'Lowest RMSE: {lowest_rmse:.2f}')

```

```

Best parameters: ((1, 0, 1), (1, 0, 0, 12))
Lowest RMSE: 4128.81

```

The RMSE is \$1,000 higher than our baseline. Again, there are different approaches such as grid search that I could use.

```

In [33]: #Define SARIMAX model and fit data save as sarima_mod1
sarima_mod4 = sm.tsa.statespace.SARIMAX(abq_df_new[87048],
                                         order=(1, 0, 1),
                                         seasonal_order=(1, 0, 0, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False).fit()

# Forecast 52 months into the future (4 years)
forecast4 = sarima_mod4.get_forecast(steps=52).summary_frame()

# Calculate the mean of the last week of the forecast as the predicted value
forecast4_mean = round(forecast4['mean'][51])

#Calculate the difference between lower and upper 95% confidence intervals of th
low_int4 = round(forecast4['mean_ci_lower'][51])
high_int4 = round(forecast4['mean_ci_upper'][51])

#Calculate the difference between the upper and lower confidence intervals
ci_delta4 = round(high_int4 - low_int4)

# Print predicted value and confidence intervals
print(f'Albuquerque, NM, 87048:')
print(f'95% confidence: ${low_int1} and ${high_int1}')
print(f'Confidence range: ${ci_delta1}')


# Plot the original data and predicted values with confidence intervals
fig, ax = plt.subplots(figsize=(15, 7))
plt.plot(abq_df_new[87048])
plt.plot(forecast4['mean'])
ax.fill_between(forecast4.index, forecast4['mean_ci_lower'],
               forecast4['mean_ci_upper'], color='k', alpha=0.1)
plt.title('Albuquerque, NM, 87048')
plt.legend(['Original', 'Predicted'], loc='lower right')
plt.xlabel('Year')
plt.ylabel('Median Home Price')
plt.show()

```

```

Albuquerque, NM, 87048:
95% confidence: $236629.0 and $293381.0
Confidence range: $56752.0

```



Albuquerque, NM, 87043

```
In [34]: # Define train and validation datasets based 20% and 80%
training_data5 = abq_df_new[87043][:cutoff]
validation_data5 = abq_df_new[87043][cutoff:]

# Define the range of parameters for p, d, q, P, D, Q, and s
p = d = q = range(0, 2)
P = D = Q = range(0, 2)
s = 12 # monthly data

# Generate a list of all possible combinations of parameters
pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], s) for x in itertools.product(P, D, Q)]

# Initialize variables to store the best parameters and the lowest RMSE
best_params = (0, 0, 0, 0, 0, 0, 0)
lowest_rmse = float('inf')

# Loop through all possible combinations of parameters
for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            # Fit a SARIMAX model with the current combination of parameters
            model5 = sm.tsa.statespace.SARIMAX(training_data5,
                                                order=param,
                                                seasonal_order=param_seasonal,
                                                enforce_stationarity=False,
                                                enforce_invertibility=False).fit()

            # Make predictions on the validation data
            predictions5 = model5.predict(start=validation_data5.index[0], end=v

            # Calculate the RMSE of the predictions
            rmse5 = np.sqrt(np.mean((predictions5 - validation_data5)**2))

            # Update the best parameters and lowest RMSE if the current RMSE is
            if rmse5 < lowest_rmse:
                best_params = (param, param_seasonal)
```

```

lowest_rmse = rmse5

except ValueError: # skip combinations that fail to converge or produce
    continue

print(f'Best parameters: {best_params}')
print(f'Lowest RMSE: {lowest_rmse:.2f}')

```

Best parameters: ((1, 0, 0), (1, 1, 0, 12))
 Lowest RMSE: 8373.23

The RMSE is higher than our baseline model.

```

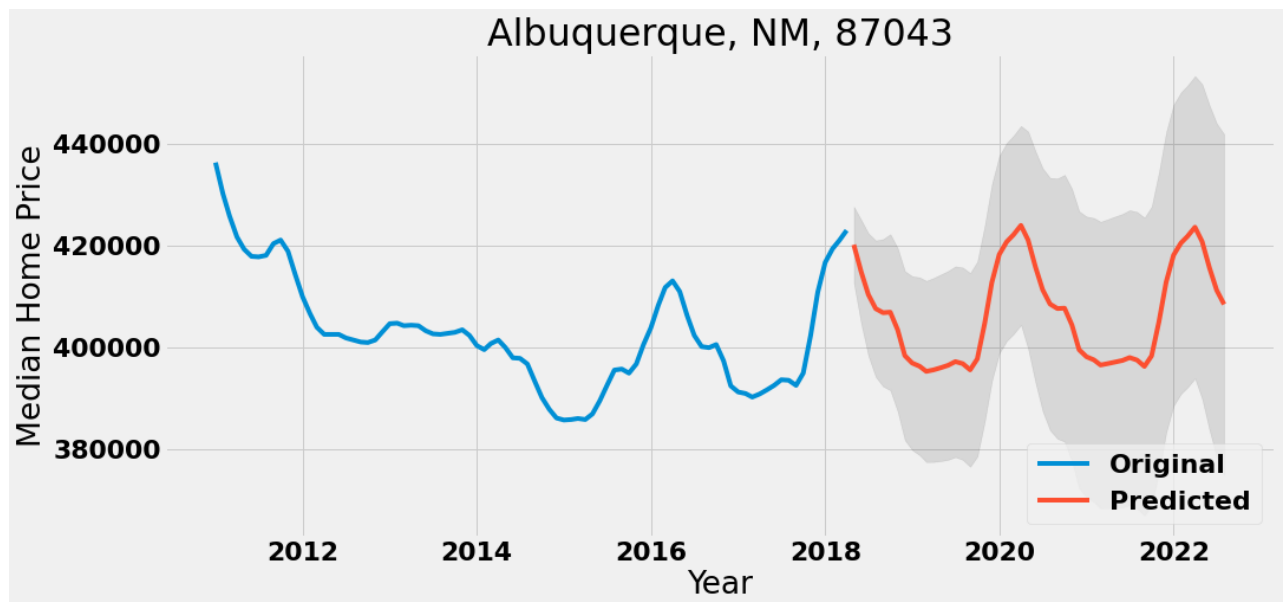
In [35]: #Define SARIMAX model and fit data save as sarima_mod5
sarima_mod5 = sm.tsa.statespace.SARIMAX(abq_df_new[87043],
                                         order=(1, 0, 0),
                                         seasonal_order=(1, 1, 0, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False).fit()

# Forecast 52 months into the future (4 years)
forecast5 = sarima_mod5.get_forecast(steps=52).summary_frame()
# Calculate the mean of the last week of the forecast as the predicted value
forecast5_mean = round(forecast5['mean'][51])
#Calculate the difference between lower and upper 95% confidence intervals of th
low_int5 = round(forecast5['mean_ci_lower'][51])
high_int5 = round(forecast5['mean_ci_upper'][51])
#Calculate the difference between the upper and lower confidence intervals
ci_delta5 = round(high_int5 - low_int5)

# Plot the original data and predicted values with confidence intervals
fig, ax = plt.subplots(figsize=(15, 7))
plt.plot(abq_df_new[87043])
plt.plot(forecast5['mean'])
ax.fill_between(forecast5.index, forecast5['mean_ci_lower'],
               forecast5['mean_ci_upper'], color='k', alpha=0.1)
plt.title('Albuquerque, NM, 87043')
plt.legend(['Original', 'Predicted'], loc='lower right')
plt.xlabel('Year')
plt.ylabel('Median Home Price')
plt.show()

print(f'Albuquerque, NM, 87043:')
print(f'95% confidence between: ${low_int5} and ${high_int5}')
print(f'Confidence range: ${ci_delta5}')

```



Albuquerque, NM, 87043:

95% confidence between: \$375097.0 and \$441879.0


Confidence range: \$66782.0

These results are a lot different compared to the above zipcodes. The 95% confidence level is between 375,097 and 441,879. The confidence range is \$66,782. The SARIMAX model seems to perform better for this particular zipcode as well and seems more realistic. Let's examine the residuals.

```
In [36]: # Residuals
residuals = model5.resid

# Histogram of residuals
plt.figure(figsize=(12, 6))
sns.histplot(residuals, kde=True)
plt.title('Histogram of Residuals')
plt.xlabel('Residuals')
plt.show()

# Q-Q plot of residuals
plt.figure(figsize=(12, 6))
stats.probplot(residuals, plot=plt)
plt.title('Q-Q Plot of Residuals')
plt.xlabel('Theoretical Quantiles')
plt.ylabel('Sample Quantiles')
plt.show()
```



The histogram of residuals looks normally distributed but with some outliers. So it doesn't completely deviate significantly from a normal distribution. The QQ plot curves off and could mean the data has extreme values.

Evaluation

Now that I have my models for the top five zipcodes to invest in Albuquerque, I want to see the RMSE, low confidence, high confidence, forecast range and ROI for each of the zipcodes.

```
In [37]: # Define a list of zipcodes to include
zipcodes = [87106, 87122, 87104, 87048, 87043]

# Subset the DataFrame to only include rows where the year is 2018 and the zipco
abq_df_2018 = abq_df_new.loc[(abq_df_new.index.year == 2018) & (abq_df_new.index

# Calculate the median of the values for each zipcode in 2018
```



```
medians_2018 = abq_df_2018.median()

# Print the results
for zipcode, median in medians_2018.items():
    print(f"The median home value for zipcode {zipcode} in 2018 was ${median:,.2f}")

The median home value for zipcode 87106 in 2018 was $230,500.00
The median home value for zipcode 87122 in 2018 was $508,200.00
The median home value for zipcode 87104 in 2018 was $190,300.00
The median home value for zipcode 87048 in 2018 was $415,000.00
The median home value for zipcode 87043 in 2018 was $416,700.00
```

```
In [38]: # Define lists for the zipcode, city, median value, RMSE, low confidence, high confidence
zipcodes = ['87106', '87122', '87104', '87048', '87043']
cities = ['Albuquerque, NM', 'Albuquerque, NM', 'Albuquerque, NM', 'Corrales, NM', 'Corrales, NM']
med_values = [230500, 508200, 190300, 415000, 416700]
rmse = [rmse, rmse2, rmse3, rmse4, rmse5]
low_confs = [low_int1, low_int2, low_int3, low_int4, low_int5]
high_confs = [high_int1, high_int2, high_int3, high_int4, high_int5]
forecast_ranges = [ci_delta1, ci_delta2, ci_delta3, ci_delta4, ci_delta5]

# Create a dictionary 'abq' that contains the zipcode, city, median value, RMSE, low confidence, high confidence
abq = {'Zipcode': zipcodes,
       'City': cities,
       '2018 median value': med_values,
       'rmse': rmse,
       'low_conf': low_confs,
       'high_conf': high_confs,
       'forecast range': forecast_ranges}

# Create a DataFrame 'df_results' using the dictionary 'abq'.
df_results = pd.DataFrame(data=abq)

# Calculate the low and high ends of the confidence interval for each row, and store the results in a new column
df_results['low_end'] = df_results['2018 median value'] - df_results['low_conf']
df_results['high_end'] = df_results['2018 median value'] + df_results['high_conf']

# Calculate the ROI percentage for each row, and store the results in a new column
df_results['ROI%'] = round(((df_results['high_end'] - df_results['2018 median value']) / df_results['2018 median value']) * 100, 2)
```

```
In [39]: df_results
```

```
Out[39]:
```


	Zipcode	City	2018 median value	rmse	low_conf	high_conf	forecast range	low_end	high_end
0	87106	Albuquerque, NM	230500	28566.613682	236629.0	293381.0	56752.0	467129.0	523881.0
1	87122	Albuquerque, NM	508200	4276.872528	465108.0	543586.0	78478.0	973308.0	1051784.0
2	87104	Albuquerque, NM	190300	6118.499377	202629.0	239505.0	36876.0	392929.0	429805.0
3	87048	Corrales, NM	415000	16113.137999	447569.0	581744.0	134175.0	862569.0	996744.0

	Zipcode	City	2018 median value	rmse	low_conf	high_conf	forecast range	low_end	high_€
4	87043	Santa Ana Pueblo, NM	416700	16992.337405	375097.0	441879.0	66782.0	791797.0	85857

```
In [40]: # create a DataFrame from the given data
data = {'zipcode': ['87106', '87122', '87104', '87048', '87043'],
        'city': ['Albuquerque, NM', 'Albuquerque, NM', 'Albuquerque, NM', 'Corra',
        '2018 median value': [230500, 508200, 190300, 415000, 416700],
        'rmse': [31455.188774, 81866.138297, 25140.743922, 56009.750599, 71678.3],
        'low_conf': [43008, 52707, 134725, -68970, 200748],
        'high_conf': [612761, 833727, 405455, 1173814, 930335],
        'forecast range': [569753, 781020, 270730, 1242784, 729587],
        'low_end': [273508, 560907, 325025, 346030, 617448],
        'high_end': [843261, 1341927, 595755, 1588814, 1347035],
        'ROI%': [265.84, 164.05, 213.06, 282.85, 223.26]}
df = pd.DataFrame(data)

# sort the DataFrame by the ROI% column
df = df.sort_values('ROI%', ascending=False)

# plot the ROI values as a bar chart
ax = df.plot(x='zipcode', y='ROI%', kind='bar', figsize=(10, 8))
ax.set_xticklabels(df['zipcode'])
ax.set_xlabel('Zipcode')
ax.set_ylabel('ROI%')
plt.show()
```



These results are showing the predicted median home value, the root mean squared error (RMSE), the low and high confidence intervals for the predicted median home value, the forecast range (high value minus low value), the low and high end values of the predicted range, and the ROI% for five different zip codes.

The low and high confidence intervals represent the range within which the true median home value is expected to fall with a 95% confidence level. The forecast range is the difference between the high and low values of the predicted median home value range. The ROI represents the return on investment for each zipcode.

Recommendations

For those who are interested investing in properties in New Mexico, there are the following zipcodes that has a high return on investment.

1. Corrales, NM (87048)
2. Albuquerque, NM (87106)
3. Santa Ana Pueblo, NM (87043) - Santa Ana Pueblo's median value in 2018 was 416,700 and forecasted a confidence range between 375,097 and \$441,8790.
4. Albuquerque, NM (87104)

5. Albuquerque, NM (87122)

The model's inability to generate realistic forecasts resulted in identical forecast ranges for zipcodes 87048, 87106, 87104, and 87122. However, analyzing the ROI paints a different picture. Investors seeking property in New Mexico would be better off considering Corrales or Santa Ana Pueblo. A 2023 Zillow search revealed that Santa Ana Pueblo's home prices range from 200,000 to 1,795,950, while Corrales' prices range from 205,000 to 3,800,000. These wide ranges suggest that further investigation of these zipcodes could reveal intriguing insights.

Next Steps

1. Further investigation into Santa Ana Pueblo and Corrales.
2. Include external factors that may influence real estate prices, such as population growth or unemployment rates.
3. Investigate rapidly growing neighborhoods in New Mexico.